

B trees

Motivation for B-trees

- Perfect BSTs are 2 trees.
- Extended to 2-3 trees.
- Extended to 2-3-4 trees.
- Can be extended to 2-3-4-5 trees, 2-3-4-5-6 trees, ...
- Can also be extended to 3-4-5 trees:
 - Each non-leaf node has 3 or 4 or 5 children; each non-leaf node has 2 or 3 or 4 keys.
 - All the leaves are at the same level.
- Can be extended to 3-4-5-6 trees, 3-4-5-6-7 trees, ...
- In general, we can have $m-(m-1)-(m-2)-\dots-(m+p)$ trees; call them (m,p) -trees.
- Perfect BSTs are $(1,1)$ -trees, 2-3 trees are $(2,1)$ -trees, 2-3-4 trees are $(2,2)$ -trees, 3-4-5-6-7 trees are $(3,4)$ -trees.

B-trees

- A special case of (m,p) trees where
 - $m = p$
 - a relaxation on the number of children/keys condition at the root
- Perfect BSTs, 2-3-4 trees, 3-4-5-6 trees are B-trees.
- A B-tree with parameter t is a search tree such that
 - (non-leaf) root node has 2 or 3 or ... or $2t$ children (1 or 2 or ... or $2t-1$ keys)
 - non-leaf nodes have t or $t+1$ or ... or $2t$ children ($t-1$ or t or ... or $2t-1$ keys)
 - all the leaf nodes are at the same level.
- 2-3-4 trees are parameter-2 B-trees conversely, not all parameter-2 B-trees are 2-3-4 trees.

Implementing B-trees

- Each node x of a parameter- t B-tree has the following fields:
 - $n(x)$: the number of keys of node x
 - $k_1(x) < k_2(x) < \dots < k_{n(x)}(x)$: the keys at node x (in increasing order)
 - $\text{leaf}(x)$: a Boolean variable to indicate whether x is a leaf node or not
 - $C_1(x), C_2(x), \dots, C_{n(x)}(x)$: pointers to children of x (NIL if x is a leaf)
- If k_i is a key at any node in the subtree $C_i(x)$ then $k_i < k_i(x) < k_{i+1}$; i.e., keys at node x separate the ranges/intervals of keys in its subtrees.
- If n, h are the number of nodes and height of a parameter- t ($t \geq 2$) B tree then

$$h \leq \log_t((n+1)/2), \quad n \geq 2t^h - 1$$

Searching in B-trees

To search for a key k at subtree rooted at node x in a B-tree:


$BTreeSearch(k, x)$

1. $i \leftarrow 1$
2. While $(i \leq n(x))$ and $(k > k_i(x))$
 $i \leftarrow i+1$
3. If $(i \leq n(x))$ and $(k = k_i(x))$ then Return(i, x)
4. If $(leaf(x) = True)$ then Return NIL
 Else Return $BTreeSearch(k, C_i(x))$

$O(\log n)$

Inserting in B-trees: top-down approach

- There may be a need to split a full node
 - A full node will have $2t-1$ (odd number of) keys.
 - The median/middle node is pushed to the parent (which will not be a full node).
 - The first $t-1$ keys become the keys of the first split child and the last $t-1$ keys become the keys of the second split child.
- As with 2-3-4 trees, it can be shown that the split in a parameter- t B-tree results in a parameter- t B-tree (may be with an increase in the height by one).
- A pictorial representation of split at a non-root node:



Splitting non-root nodes in B-trees

To split a non-root node x in a parameter- t B-tree with p as $\text{Parent}[x]$, x as the i -th child of p ; i.e., $x = C_i[p]$, z as the newly added child of p ; i.e., z is the $i+1$ -th child of p after the split (and is the second half of x before it was split): $O(t) = O(1)$

$\text{BTreeSplit}(p, i, x)$

1. $z \leftarrow \text{NewBTreeNode}()$

2. $\text{leaf}(z) \leftarrow \text{leaf}(x)$ Initializing the fields for new node z

3. $n(z) \leftarrow t-1$

4. for $j=1$ to $t-1$
 $k_j(z) \leftarrow k_{j+t}(x)$ Assigning keys for z

5. if $\text{leaf}(x) = \text{false}$ then

 for $j=1$ to t Assigning children for z

$C_j(z) \leftarrow C_{j+t}(x)$

6. $n(x) \leftarrow t-1$ Updating the number of keys at x

7. for $j=n(p)$ down to $i+1$
 $C_{j+1}(p) \leftarrow C_j(p)$ Updating/shifting the right-children of p

8. $C_{i+1}(p) \leftarrow z$ Making z as the i -th child of p

9. for $j=n(p)$ down to $i+1$
 $k_{j+1}(p) \leftarrow k_j(p)$ Updating/shifting the right-keys of p

10. $k_i(p) \leftarrow k_i(x)$ Pushing the key from (old) x to p

11. $n(p) \leftarrow n(p)+1$ Updating the number of keys at p

Insertion into non-full nodes in B-trees

To insert key k into a non-full node x in a parameter- t B-tree

$BTreeInsertNonFull(x, k)$

1. $i \leftarrow n(x)$

2. If $leaf(x) = \text{True}$ then

 while $(i \geq 1) \text{ and } (k < k_i(x))$

$k_{i+1}(x) \leftarrow k_i(x)$

$i \leftarrow i - 1$

$k_i(x) \leftarrow k$

$n(x) \leftarrow n(x) + 1$

Else

 while $(i \geq 1) \text{ and } (k < k_i(x))$

$i \leftarrow i - 1$

$i \leftarrow i + 1$

 If $(n(C_i(x)) = 2t - 1)$ then

$BTreeSplit(x, i, C_i(x))$

 If $k > k_i(x)$ then $i \leftarrow i + 1$

$BTreeInsertNonFull(C_i(x), k)$

Run-time for a single split is $O(t)$, during insertion, there may be up to h splits; so total run-time is $O(t h) = (t \log_t n) = O(\log n)$

Insertion into in B-trees

To insert key k into a parameter- t B-tree T (by subsuming the case of splitting at the root)

$BTreeInsert(T, k)$

1. $r \leftarrow \text{root}(T)$

2. If $n(r) = 2t - 1$ then

$s \leftarrow \text{NewBTreeNode}()$

$\text{root}(T) \leftarrow s$

$\text{leaf}(s) \leftarrow \text{FALSE}$

$n(s) \leftarrow 0$

$C_1(s) \leftarrow r$

$BTreeSplit(s, 1, r)$

$BTreeInsertNonFull(s, k)$

Else

$BTreeInsertNonFull(r, k)$


Inserting in a B-tree

Construct a parameter-3 B-tree for the set {7, 29, 15, 3, 17, 11, 23, 37, 27, 13, 49, 25, 43, 55, 9, 33, 39, 51, 53, 41, 10, 20}


Note: Since the parameter for the B-tree is 3,

- Each non-root node should have 2 or 3 or 4 or 5 keys
- Root node should have 1 or 2 or 3 or 4 or 5 keys
- A full node will have 5 keys in it.


Initial tree (empty tree):



Insert 7:



Insert 29, 15, 3, 17:

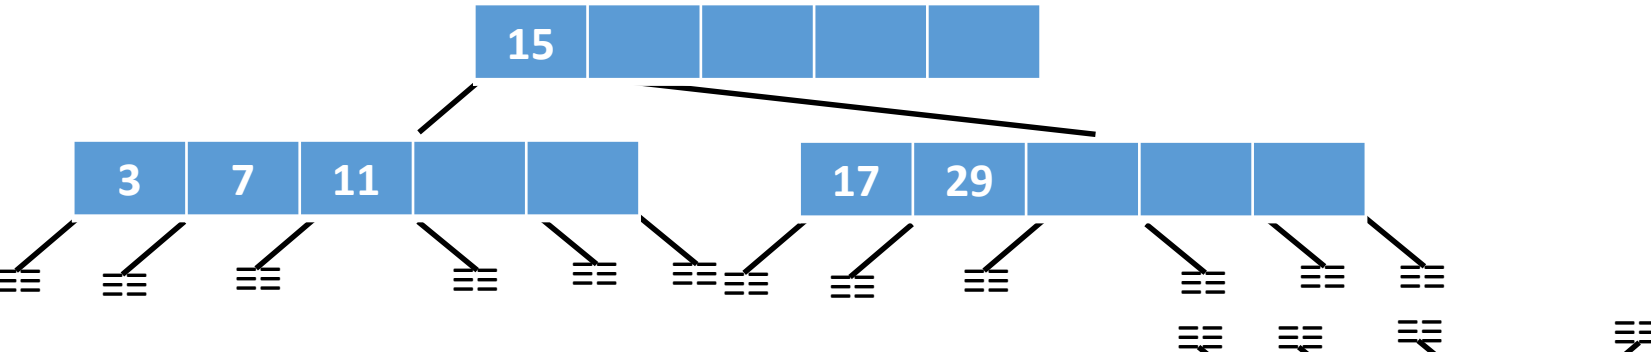



7	29	29	29	29
3	15	15	17	
	7			


Insert 11: Can't enter the root because it is a full node, so split before continuing with the search. The median at the full node becomes a key of the (new) root node.

Inserting in a B-tree

Construct a parameter-3 B-tree for the set {7, 29, 15, 3, 17, 11, 23, 37, 27, 13, 49, 25, 43, 55, 9, 33, 39, 51, 53, 41, 10, 20}




Insert 23, 37, 27, 13:





Inserting in a B-tree


Construct a parameter-3 B-tree for the set {7, 29, 15, 3, 17, 11, 23, 37, 27, 13, 49, 25, 43, 55, 9, 33, 39, 51, 53, 41, 10, 20}



Insert 49: The node [17, 23, 27, 29, 37] is on the search path and is full, so split it before continuing with the search




Insert 25, 43, 55, 9, 33:





Error: 49 is missing henceforth

Inserting in a B-tree


Construct a parameter-3 B-tree for the set {7, 29, 15, 3, 17, 11, 23, 37, 27, 13, 49, 25, 43, 55, 9, 33, 39, 51, 53, 41, 10, 20}



Insert 39: The node [29, 33, 37, 43, 55] is on the search path and is full, so split it before continuing with the search




Insert 51, 53:





Inserting in a B-tree

Construct a parameter-3 B-tree for the set {7, 29, 15, 3, 17, 11, 23, 37, 27, 13, 49, 25, 43, 55, 9, 33, 39, 51, 53, 41, 10, 20}




Insert 41: The node [39, 43, 51, 53, 55] is on the search path and is full, so split it before continuing with the search





Inserting in a B-tree

Construct a parameter-3 B-tree for the set {7, 29, 15, 3, 17, 11, 23, 37, 27, 13, 49, 25, 43, 55, 9, 33, 39, 51, 53, 41, 10, 20}




Insert 10: The node [3, 7, 9, 11, 13] is on the search path and is full, so split it before continuing with the search





Inserting in a B-tree

Construct a parameter-3 B-tree for the set {7, 29, 15, 3, 17, 11, 23, 37, 27, 13, 49, 25, 43, 55, 9, 33, 39, 51, 53, 41, 10, 20}



Insert 20: The node [9, 15, 27, 37, 51] is on the search path and is full, so split it before continuing with the search



Inserting in a B-tree

Construct a parameter-3 B-tree for the set $\{7, 29, 15, 3, 17, 11, 23, 37, 27, 13, 25, 43, 55, 9, 33, 39, 51, 53, 41, 10, 20\}$ with the insertion order being

1. $\langle 3, 7, 9, 10, 11, 13, 15, 17, 20, 23, 25, 27, 29, 33, 37, 39, 41, 43, 51, 53, 55 \rangle$
2. $\langle 55, 53, 51, 43, 41, 39, 37, 33, 29, 27, 25, 23, 20, 17, 15, 13, 11, 10, 9, 7, 3 \rangle$
3. $\langle 3, 7, 10, 11, 13, 17, 20, 23, 25, 29, 33, 39, 41, 43, 53, 55, 9, 15, 37, 51, 27 \rangle$
4. $\langle 27, 9, 15, 37, 51, 3, 7, 10, 11, 13, 17, 20, 23, 25, 29, 33, 39, 41, 43, 53, 55 \rangle$

How structurally similar are these trees?

Deleting from a B-tree

Let T be a parameter- t B-tree. To delete key k from T .

Top-down approach

- Search (recursively) for the location of the key by starting the search at the root of T .
- While searching, never enter a minimal node; i.e., a node with exactly $t-1$ keys in it (similar to not entering a full node during insertion).
 - “redistribute” before continuing the search at the minimal node (similar to splitting during insertion).
- Delete k from its location by distinguishing the following cases:
 - k is in a (non-minimal) leaf node.
 - k is in a non-leaf node which further distinguishes into cases:
 - root of the predecessor subtree is not minimal
 - root of the successor subtree is not minimal
 - root of both these subtrees are minimal

Redistribution during deletion in a B-tree

Let T be a parameter- t B-tree. To delete key k from T .

Assume that

- the search path has reached node x
- the child $C_i(x)$ is the next node on the search path but is minimal; i.e., $n(C_i(x)) = t-1$.


Case 1: (Immediate) Left or right sibling of $C_i(x)$ is non-minimal; i.e., $n(C_{i-1}(x)) > t-1$ or $n(C_{i+1}(x)) > t-1$.

Case 2: Both the (immediate) left and right siblings of $C_i(x)$ are minimal; i.e., $n(C_{i-1}(x)) = t-1$ and $n(C_{i+1}(x)) = t-1$.

Redistribution when a sibling in non-minimal

Case 1: (Immediate) Left or right sibling of $C_i(x)$ is non-minimal,
w.l.o.g, let the right sibling; i.e., $C_{i+1}(x)$ be non-minimal

- Make $C_i(x)$ non-minimal by
 - moving a key from x and
 - moving a key from non-minimal sibling to x and
 - moving a child of the non-minimal sibling as a child of $C_i(x)$




$$n(x) \geq t-1$$

$$n(C_i(x)) = t-1$$

$$n(C_{i+1}(x)) \geq t$$

$$T_{1,1} < \epsilon < T_{1,2} < \alpha < T_{1,3} < \beta \text{ \& } \beta < T_{2,1} < \gamma < T_{2,2} < \delta < T_{2,3}$$



$$n(x) \geq t-1+1-1;$$

lost β , got γ

$$n(C_i(x)) = t-1+1;$$

got β


$$n(C_{i+1}(x)) \geq t-1;$$

lost γ

$$T_{1,1} < \epsilon < T_{1,2} < \alpha < T_{1,3} < \beta < T_{2,1} < \gamma \text{ \& } \gamma < T_{2,2} < \delta < T_{2,3}$$


Redistribution when a sibling in non-minimal

Example: Delete 5 from the following parameter-3 B-tree.



Here, $k=5$

The node to be entered on the search path is $[3, 5]$ which is minimal. It has a non-minimal (right) sibling. So, we do the redistribution with $x=[7, 25, 33, 43, 51]$, $C_i(x)=[3, 5]$, $C_{i+1}(x)=[11, 21, 23]$.




Redistribution when both the siblings are minimal

Case 2: Both the (Immediate) left and right sibling of $C_i(x)$ are minimal,

- Make $C_i(x)$ non-minimal by
 - merging $C_i(x)$ with one of its siblings, say its right sibling $C_{i+1}(x)$
 - move a key from x to the (new) merged node such that this key becomes the median of the merged node.

Before




$$n(x) \geq t$$

$$n(C_i(x)) = t-1$$

$$n(C_{i+1}(x)) = t-1$$

After



$$n(x) \geq t-1;$$

lost β


$$n(C_i(x)) = t-1+1+t-1 = 2t-1;$$

got β and all the keys from $C_{i+1}(x)$

$C_{i+1}(x)$ DNE; lost all its keys


Redistribution when both the siblings are minimal

Example: Delete 9 from the following parameter-3 B-tree.



Here, $k=9$

The node to be entered on the search path is [33] which is minimal. It is a root node. We do the redistribution with $x=[33]$, $C_i(x)=[7, 25]$, $C_{i+1}(x)=[43, 51]$.



Deleting from a B-tree

Let T be a parameter- t B-tree. To delete key k from T .

- Let N be the node where k is located,
- having ensured that the search path never entered a minimal node,
- though after redistribution some nodes on the search path (after passing through them) may have become minimal.
- Cases to be considered:
 - Case 1: N is a leaf node
 - Case 2: N is a non-leaf node
 - Case 2.1: Root of the predecessor (w.r.t. k) subtree of N is non-minimal
 - Case 2.2: Root of the predecessor subtree of N is minimal while the successor subtree is non-minimal
 - Case 2.3: Both the predecessor and successor subtrees of N are minimal


Deleting from a B-tree

Case 1: N is a leaf node

Note: N is non-minimal; otherwise the search path would not have entered N. So,

- *remove the key k from N and*
- *reduce $n(N)$ by one*

Example: To delete 13 from the following B(3)-tree (assuming that redistribution has been done)




Deleting from a B-tree

Case 2.1: N is a non-leaf node, the root M of the predecessor subtree is non-minimal.

- traverse along the subtree rooted at M to locate the replacement key $k' =$ predecessor of k ; let P be the node containing k' [How to identify the predecessor key?]
- delete (recursively) the replacement key, i.e.,
 - at N , replace k with k' and at P , delete k'

Example: To delete 27 from the following $B(3)$ -tree (assuming redistribution)

Here, $k=27$, $N=[7, 15, 27]$, $M=[21, 23, 25]$, $k'=25$, $P=M=[21, 23, 25]$




Deleting from a B-tree

Case 2.2: N is a non-leaf node, the root O of the predecessor subtree is non-minimal.

- traverse along the subtree rooted at O to locate the replacement key $k' =$ successor of k ; let P be the node containing k' [How to identify the successor key?]
- delete (recursively) the replacement key, i.e.,
 - at N , replace k with k' and at P , delete k'

Example: To delete 15 from the following $B(3)$ -tree (assuming redistribution)

Here, $k=15$, $N=[7, 15, 27]$, $O=[21, 23, 25]$, $k'=21$, $O=M=[21, 23, 25]$



Deleting from a B-tree

Case 2.3: N is a non-leaf node, the root M of the predecessor subtree and the root O of the successor subtree are minimal.


- merge nodes M and O by moving k into this merged node (let M be the merged node, release node O , shift children if require).

[Now, M is full (non-minimal), k is the median of M]

- delete (recursively) k from M .

Example: To delete 15 from the following B(3)-tree (assuming redistribution)

Here, $k=15$, $N=[7, 15, 27]$, $M=[9, 11]$, $O=[21, 23]$



Exercise

Write the pseudocode and compute the run-time for the following operations on a $B(t)$ -tree:

- Redistribution
- Identifying predecessor key
- Identifying successor key
- Deleting a key